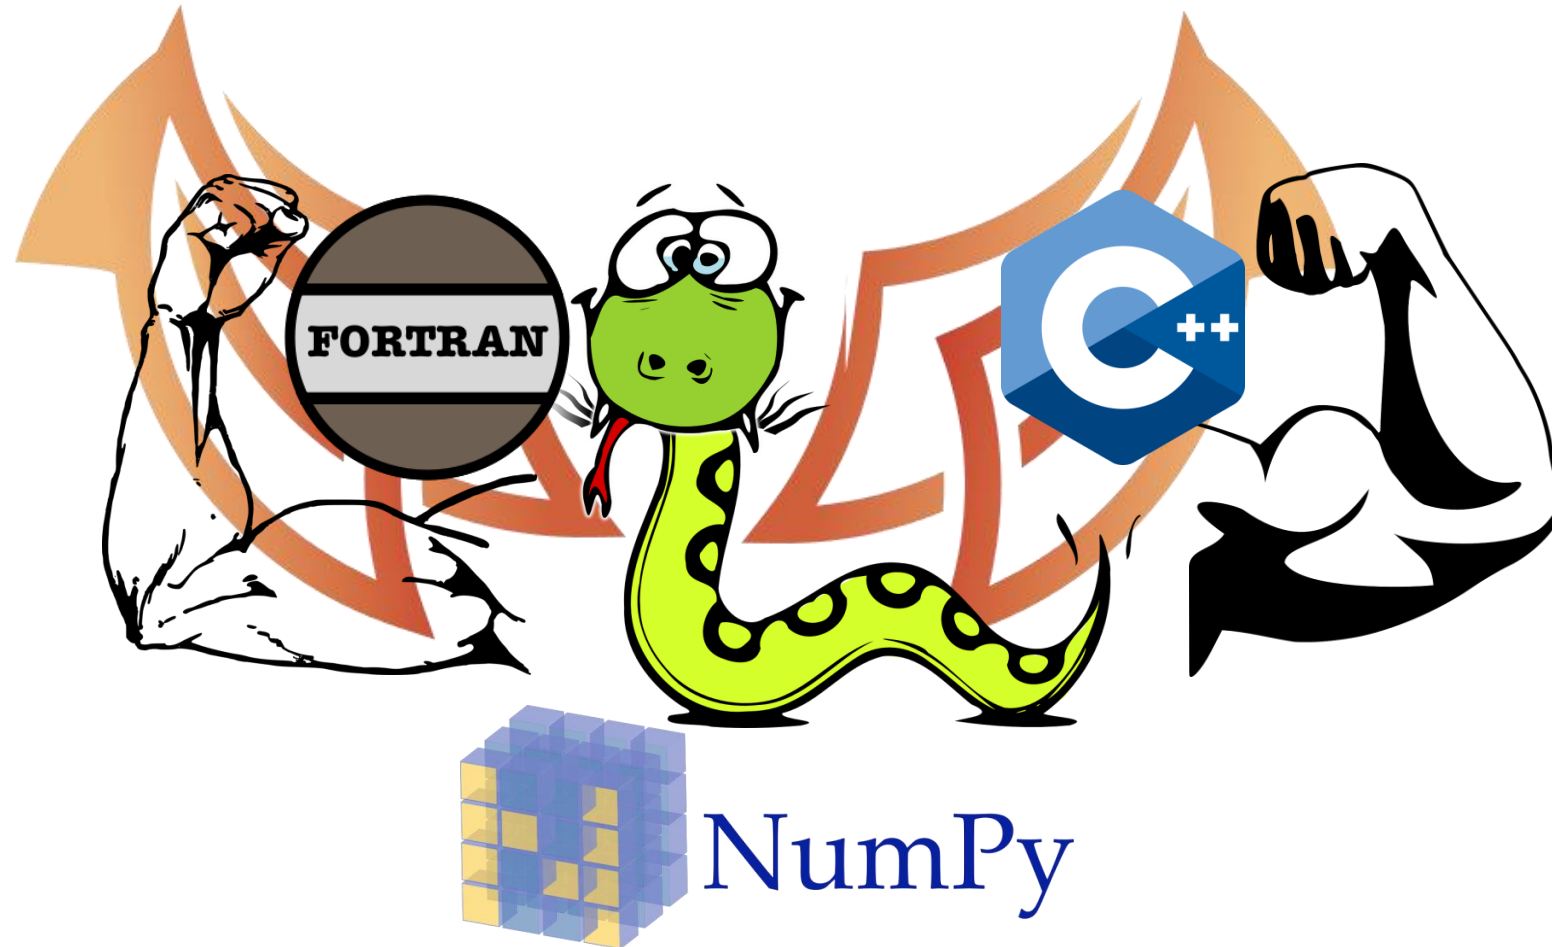


# Parallelize Python using Dask

Yubo “Paul” Yang, THW-IL, 2019/11/20

Code available on [GitHub](https://github.com/Paul-St-Young/thw-dask-para): <https://github.com/Paul-St-Young/thw-dask-para>

Images from <https://pixabay.com/vectors>

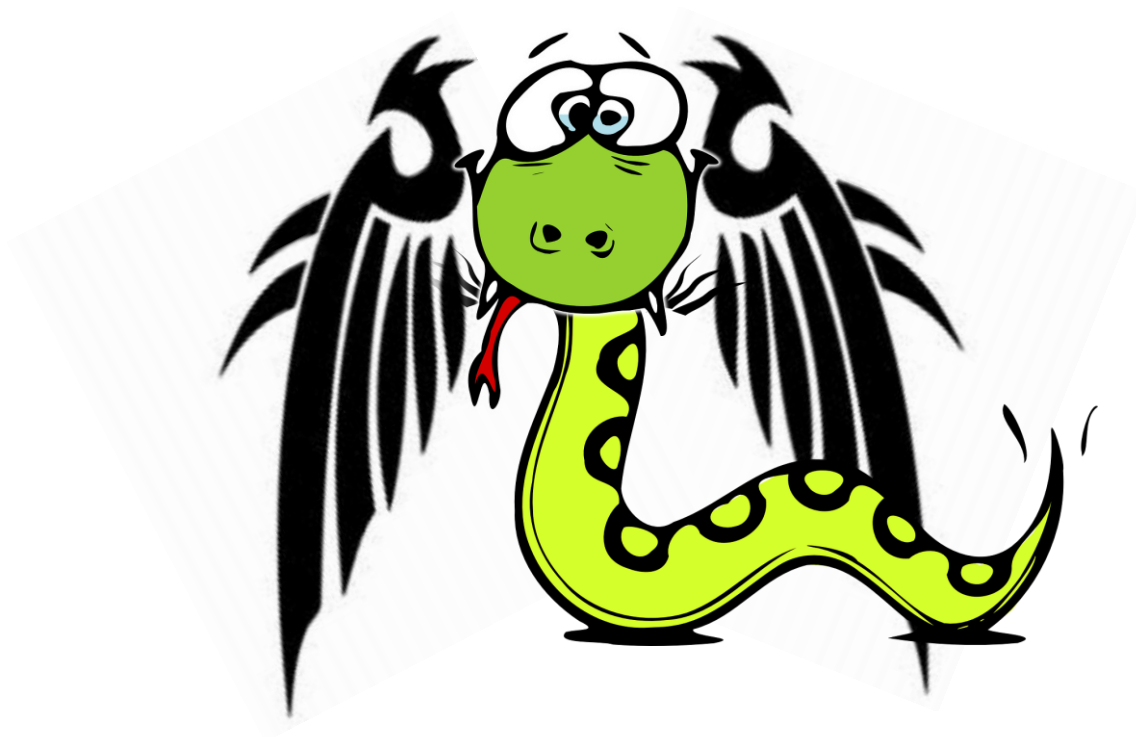


# Why parallelizing Python is problematic: The Global Interpreter Lock (**GIL**)

The GIL **prevents multiple threads** from executing CPython bytecodes at once.

This lock is necessary mainly because CPython's **memory management is not thread-safe**.

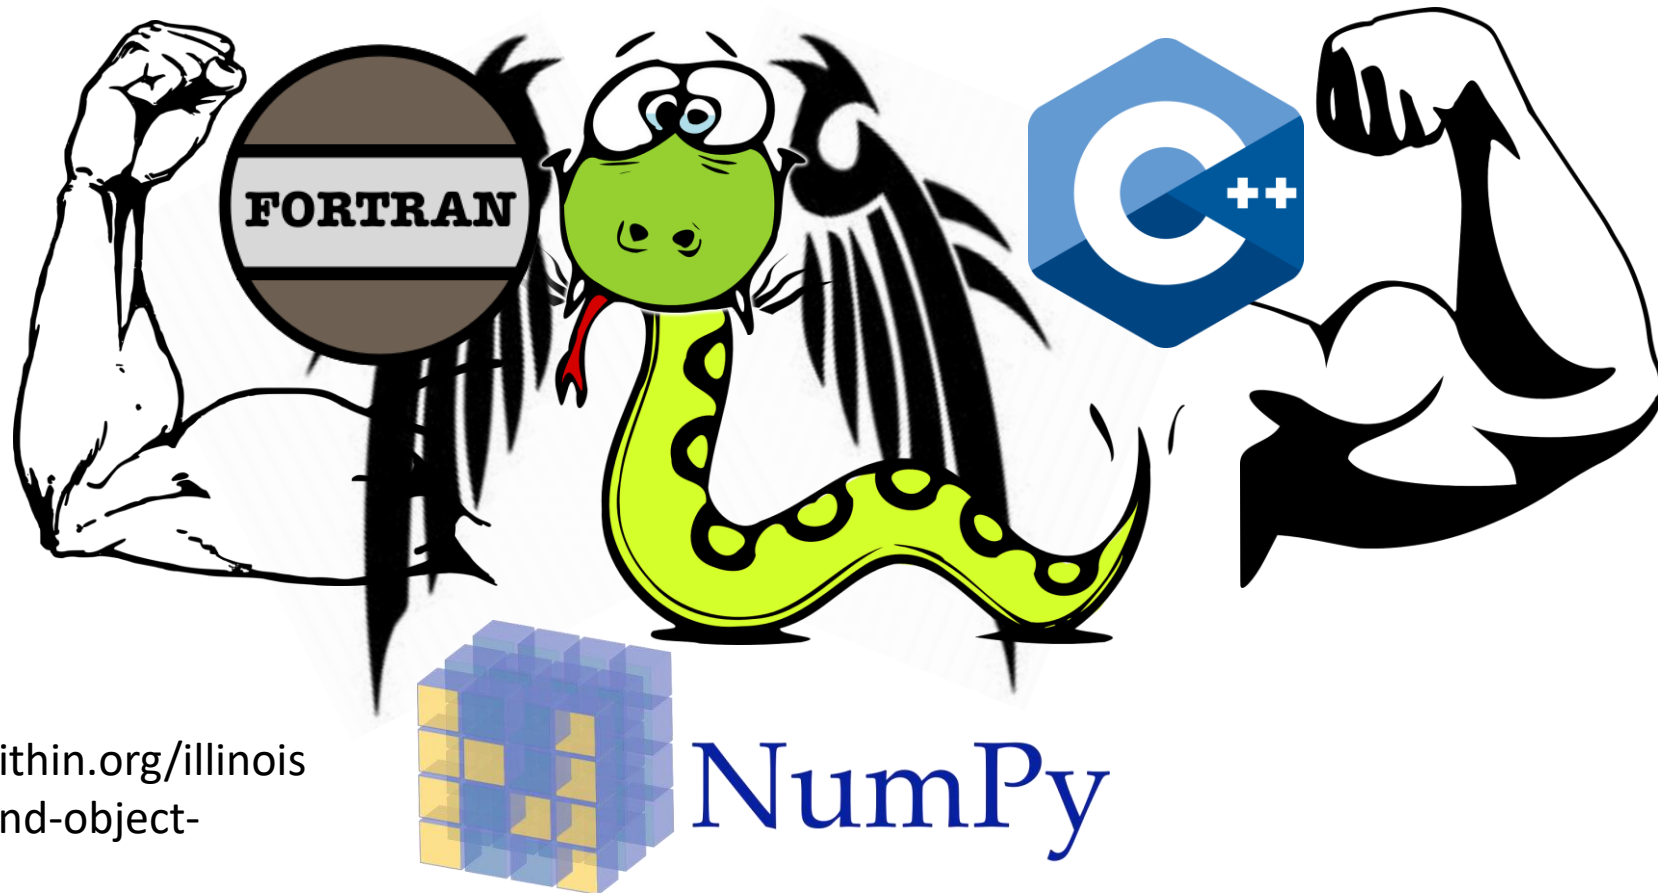
Take home: pure Python script is dragged down by the GIL to a single processing thread!



# Why parallelizing Python is problematic: The Global Interpreter Lock (**GIL**)

The GIL can be circumvented by calling low-level C/Fortran functions, which numpy does!

*Problem:* require significant coding, and still limited to a single machine.



See 2019/04/17 [talk](http://www.thehackerwithin.org/illinois/posts/python-as-glue-and-object-oriented-programming)  
<http://www.thehackerwithin.org/illinois/posts/python-as-glue-and-object-oriented-programming>

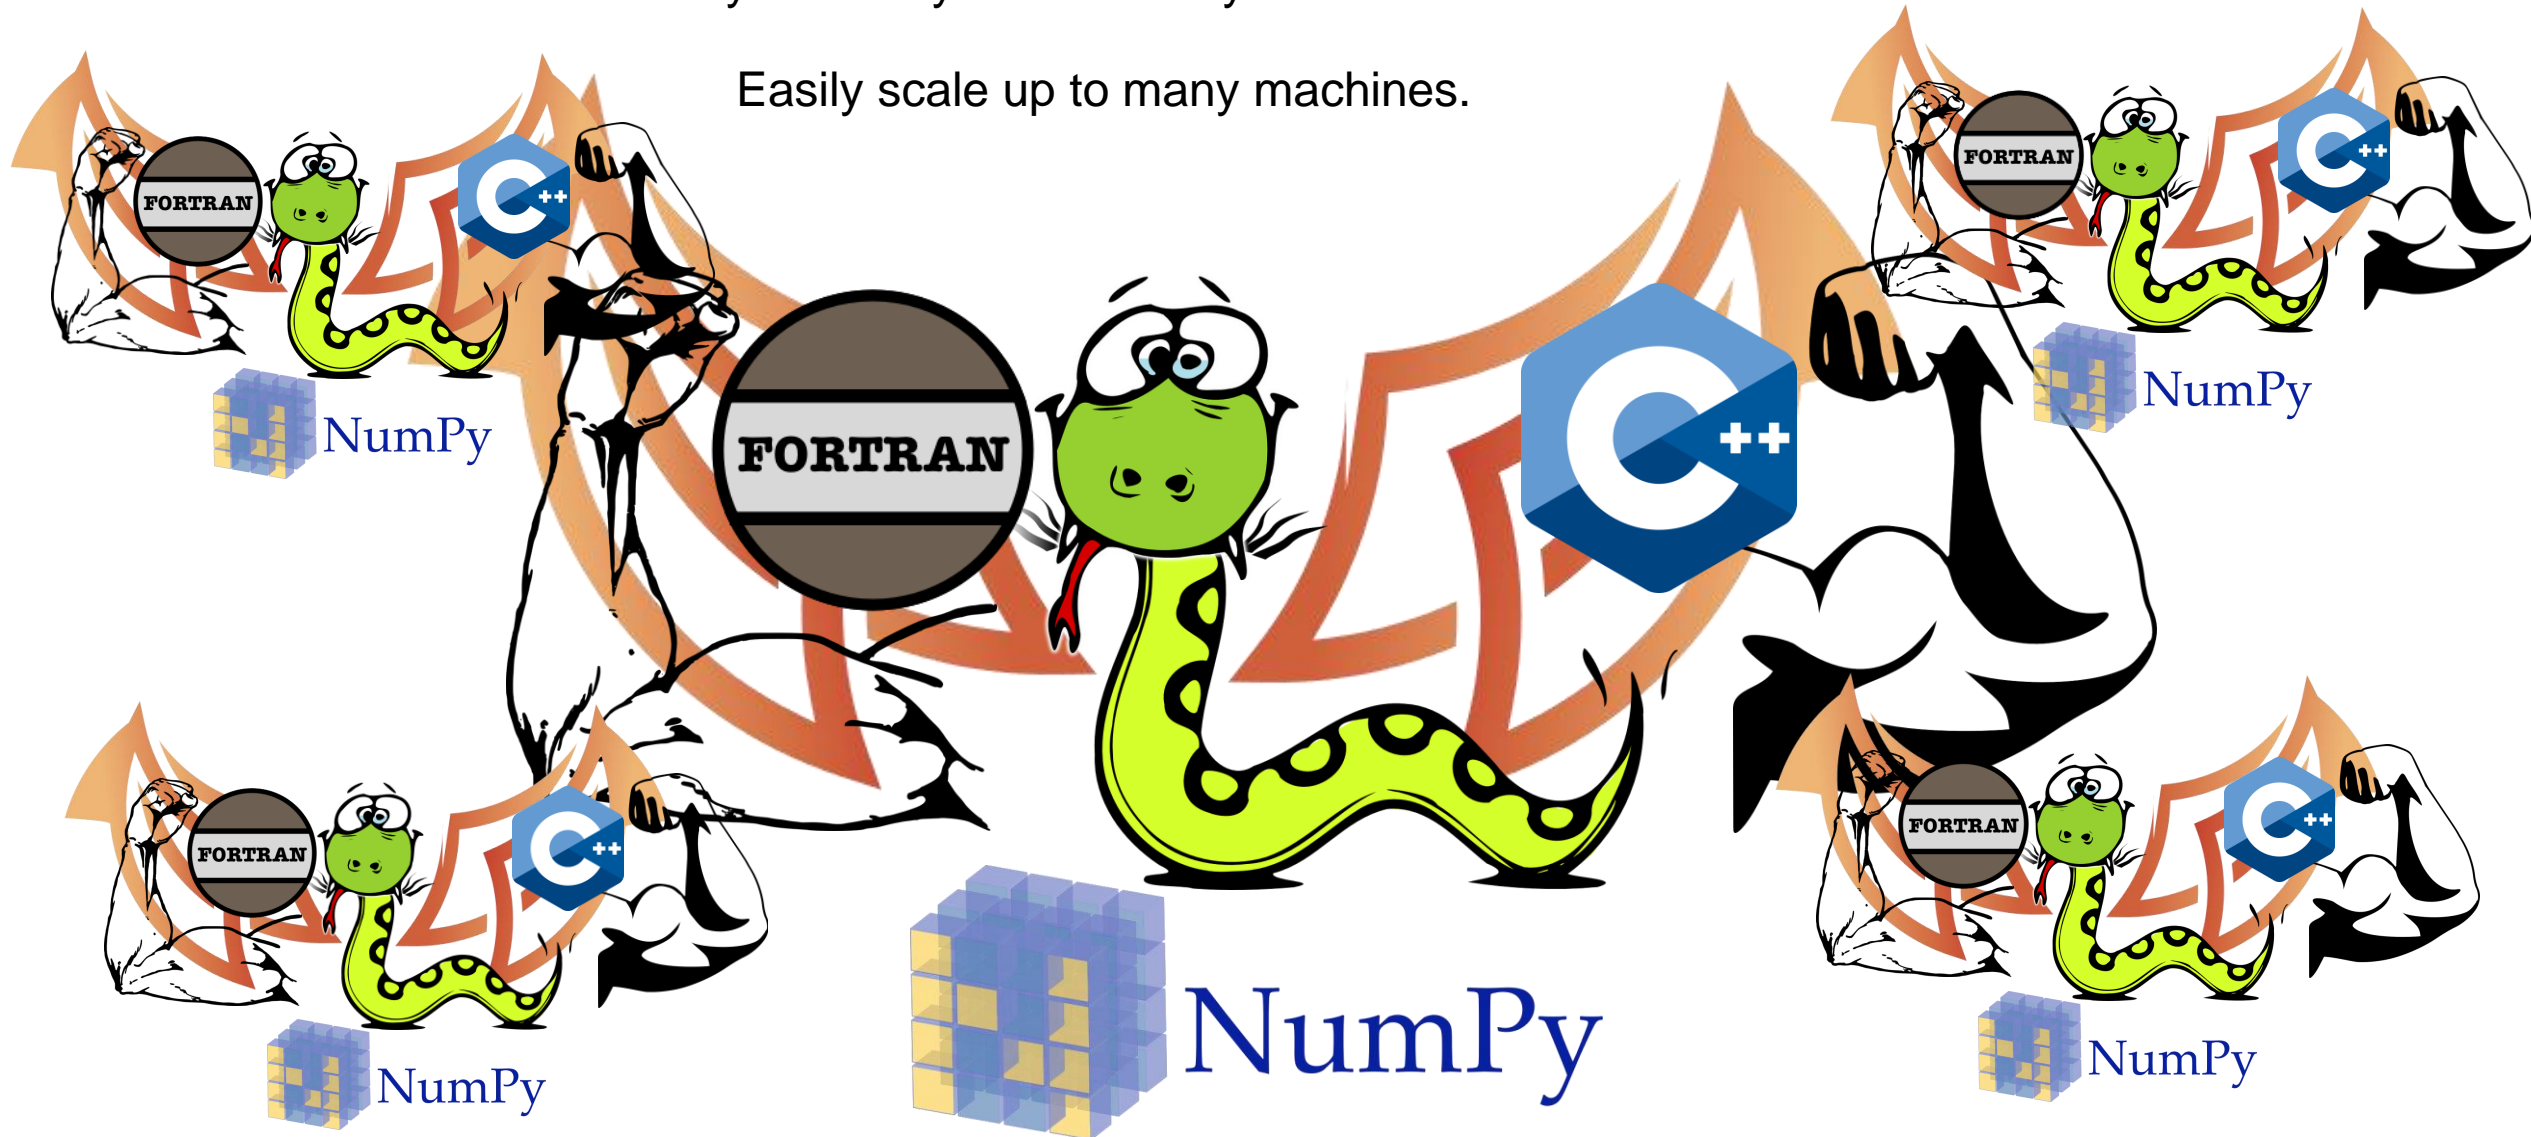
# What is Dask? Graph execution engine

Dask is a graph execution engine.

Dask is the most intuitive way to run Python on many cores.

dask uses: psutil, multiprocessing, mpi4py, and sockets under the hood.

Easily scale up to many machines.



# Why Dask?

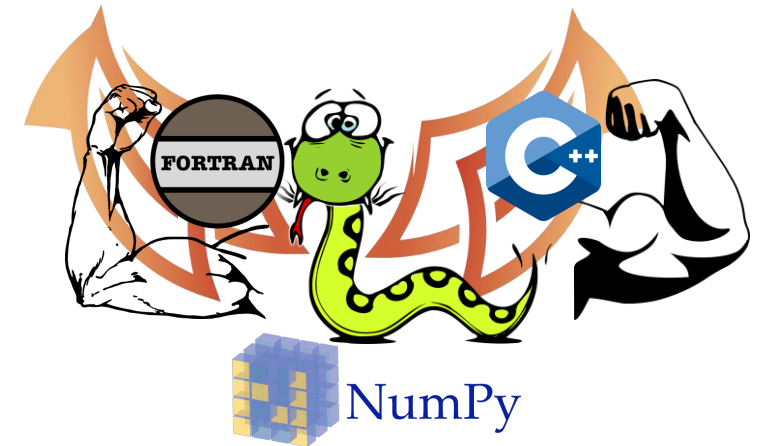
*Reason 1: Dask stands on solid foundation*

- “Dask does not seek to disrupt or displace the existing ecosystem, but rather to complement and benefit it from within.”
- strong and diverse backers

*Reason 2: Dask developers care*

- well-written documentation: docs, distributed, ml, [examples.dask.org](https://examples.dask.org)
- easy to diagnose: dashboard, serial->supercomputer
- many video tutorials: [showcase](#), [dashboard](#), [lab extension](#)

*Reason 3: Dask reaches the cloud*





# How to use Dask? Complete pipeline $\equiv$ create task graph $\rightarrow$ scheduler $\rightarrow$ execute

Start by making one choice for each step:

task graph builder: `dask.delayed`

task graph executioner: `dask.distributed.Client(LocalCluster())`

**Collections**  
(create task graphs)

**Task Graph**

**Schedulers**  
(execute task graphs)

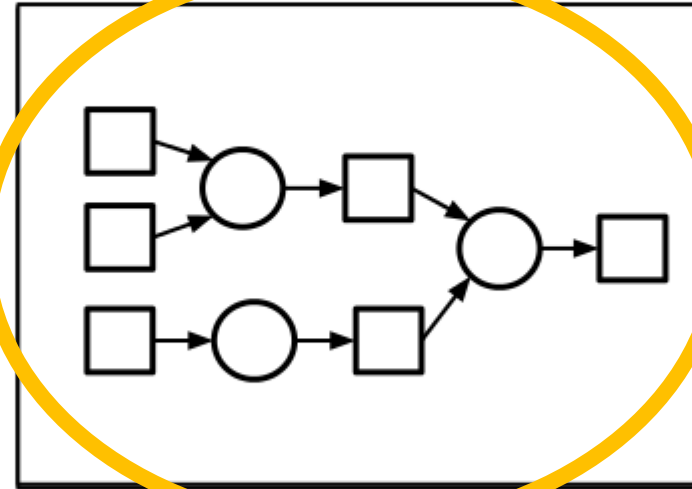
Dask Array

Dask DataFrame

Dask Bag

Dask Delayed

Futures



Single-machine  
(threads, processes,  
synchronous)

Distributed

## USER INTERFACE

User Interfaces

Array

Bag

DataFrame

Delayed

Futures

Machine Learning

Best Practices

API

## SCHEDULING

Scheduling

Distributed Scheduling

## DIAGNOSTICS

Understanding Performance

Visualize task graphs

Diagnostics (local)

Diagnostics (distributed)

# Step 1: Build and visualize the task graph

Use `dask.delayed` to lazy-evaluate costly computation. This constructs the task graph without doing any real computation.

A delayed object can be visualized

or executed.

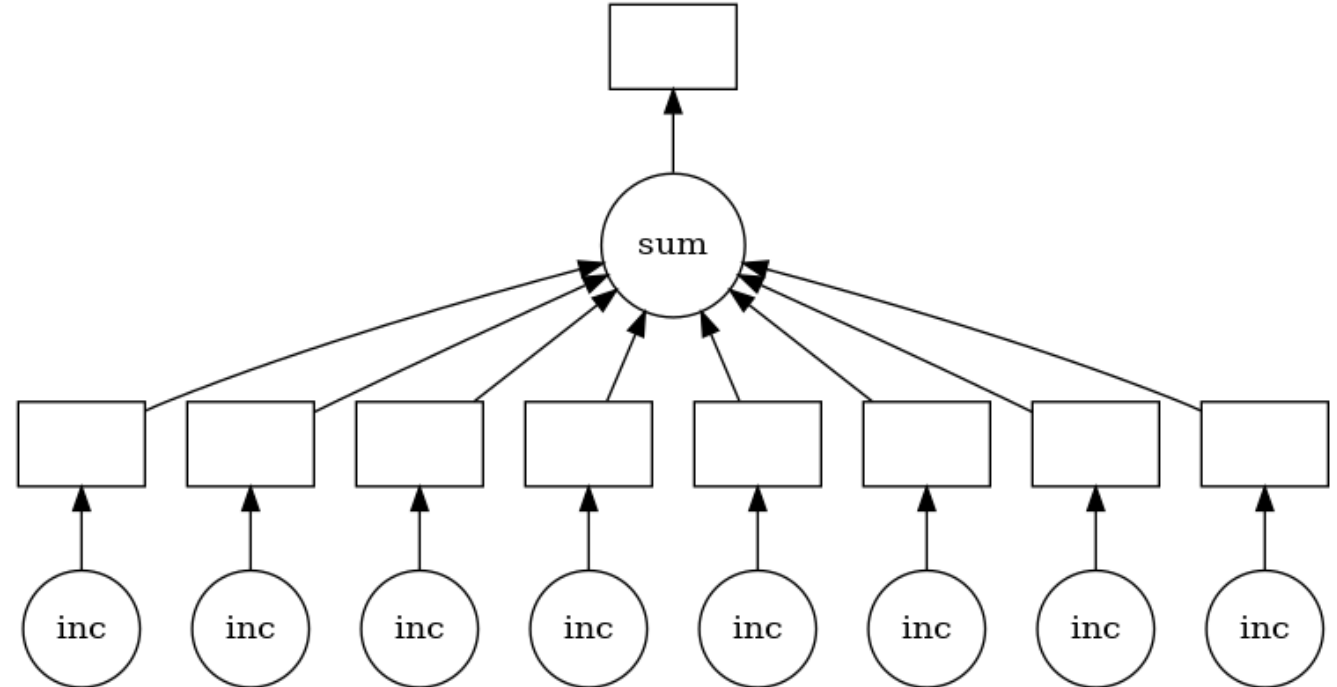
```
%%time
results = []

for x in data:
    y = delayed(inc)(x)
    results.append(y)

total = delayed(sum)(results)
print("Before computing:", total) # Let's see what type of thing total is
result = total.compute()
print("After computing :", result) # After it's computed
```

```
Before computing: Delayed('sum-f61ea7b0-b27f-4e1b-9507-3d0154fb6840')
After computing : 44
CPU times: user 14.9 ms, sys: 23.1 ms, total: 37.9 ms
Wall time: 1.03 s
```

```
total.visualize()
```

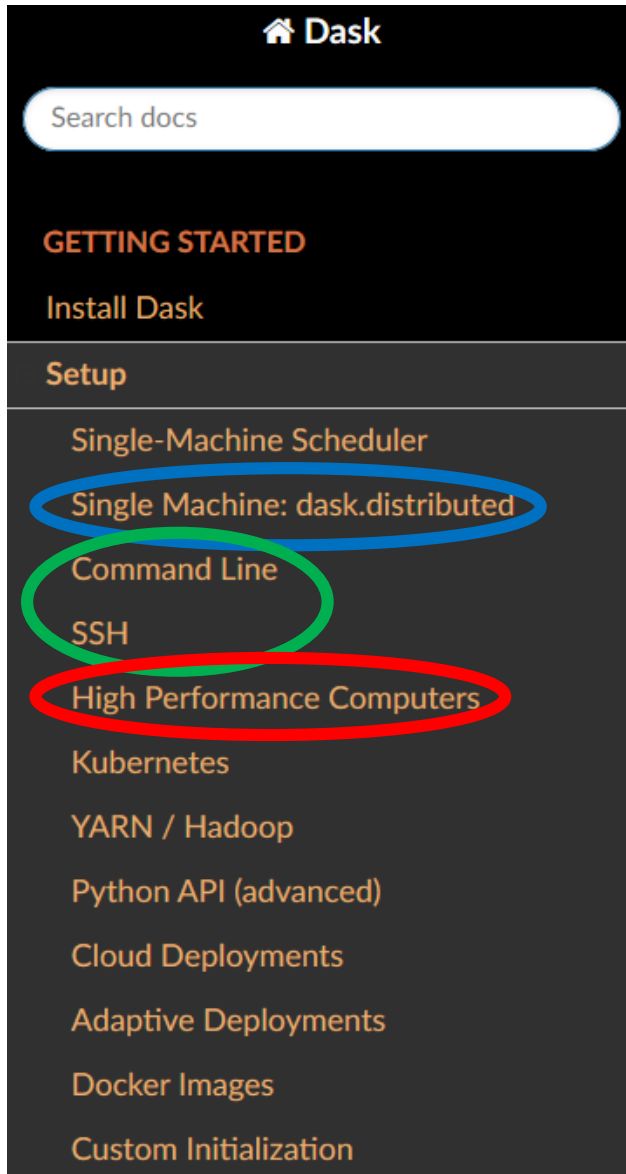


Example taken from dask tutorial

[Tutorial 01](#) delayed task graph

[https://github.com/dask/dask-tutorial/blob/master/01\\_dask.delayed.ipynb](https://github.com/dask/dask-tutorial/blob/master/01_dask.delayed.ipynb)

## Step 2: Construct a scheduler to efficiently **execute** the task graph



**LocalCluster** processes=False -> like pure OpenMP, one process saves memory and communication, but Python script must release the GIL!

**LocalCluster** processes=ncpu -> like pure MPI, no need to worry about GIL, but uses ncpu × memory

**Client**('127.0.0.1:8787') -> use a few loosely connected workstations to pool memory and CPUs together

Requires setup:

on local machine:

dask-scheduler

dask-worker [tcp://myscheduler:8787] --nprocs 8

on remote machine:

dask-worker [tcp://myscheduler:8787] --nprocs 8

**PBSCluster** -> use a super computer with torque scheduler  
cluster = PBSCluster(cores=16, memory="16GB", queue="secondary",  
walltime="00:15:00", interface="ib0")  
cluster.scale(jobs=4) # get 4 nodes

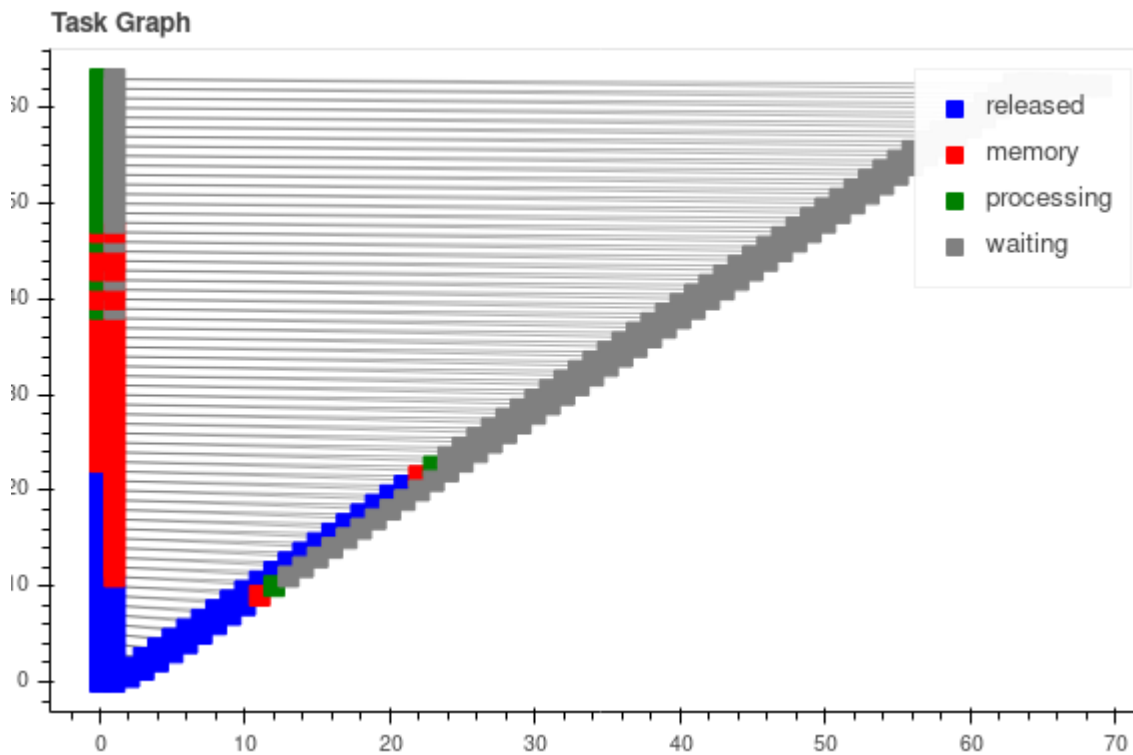


# Step 3: Use dashboard to profile and optimize

One task stream for each worker thread,  
want to see lots of solid colors (computation).

Progress bar show headroom for more cores.

Task graph shows parallelism  
inherit in computation.



# Conclusions

Dask is an intuitive and flexible graph execution engine for Python.

Dask.delayed can easily parallelize existing Python code.

Dask.distributed.Client can use a variety of Clusters to suit the task at hand.

Eg. Laptop -> LocalCluster,

Workstations -> SSHCluster,

Supercomputer -> PBSCluster\*

Cloud->KubeCluster

\*note: use dask.future instead of dask.delayed if Client has a job submission system

Dask dashboard provides many tools to debug and profile parallel execution.

# Resources

Videos by Matthew Rocklin

[cmd setup](#) cluster  
[dashboard](#)  
[delayed](#)

Tutorials

[dask/dask-tutorial](#)

Documentation

**examples.dask.org**  
docs.dask.org  
distributed.dask.org

array, bag, database APIs

**Take-home:** smash that GIL with f2py/numpy/pybind, then dask to clusters!

